

# NINAD SAMUDRE

CLOUD ENGINEER AND TRAINER

## Docker and Kubernetes



CGI



# Training Objectives

---



At the end of training, participants should be able to

- Know Docker & swim with them
- Bundle applications in Docker images
- Run Docker Containers



# Table of Contents



## Module 1: Docker concepts and terms

- Intro
- Containerization vs Virtualization
- Terminologies in Docker world
- Docker Architecture
- Docker Machine
- Configuring Docker engine

## Module 2: Docker Containers

- Creating containers
- Running containers
- Running containers in background
- Connecting containers
- Docker Images

## Module 3: Provisioning Docker Image

- Introducing the Dockerfile
- Creating a Dockerfile
- Building images manually
- Building images using Continuous Integration tools
- Storing and retrieving Docker Images from Docker Hub
- Inspecting a Dockerfile from DockerHub

## Module 4: Diving Deeper into Dockerfile

- The Build cache
- Dockerfile and Layers
- Building a Web Server Container
- The CMD Instruction Docker
- The ENTRYPOINT Instruction
- The ENV Instruction
- Volumes and the VOLUME Instruction

## Module 5: Working with Registry

- Module Intro
- Creating a Public repo on Docker Hub
- Using our Public repo on Docker Hub,
- Using a Private Registry,
- Docker Hub Enterprise

## Module 6: Docker Networking

- Module Intro
- The docker0 Bridge
- Virtual Ethernet Interfaces
- Network Configuration Files
- Exposing Ports
- Viewing Exposed Ports
- Linking Containers
- Lab Exercises



# Module 1: Docker Concept & Terms

---



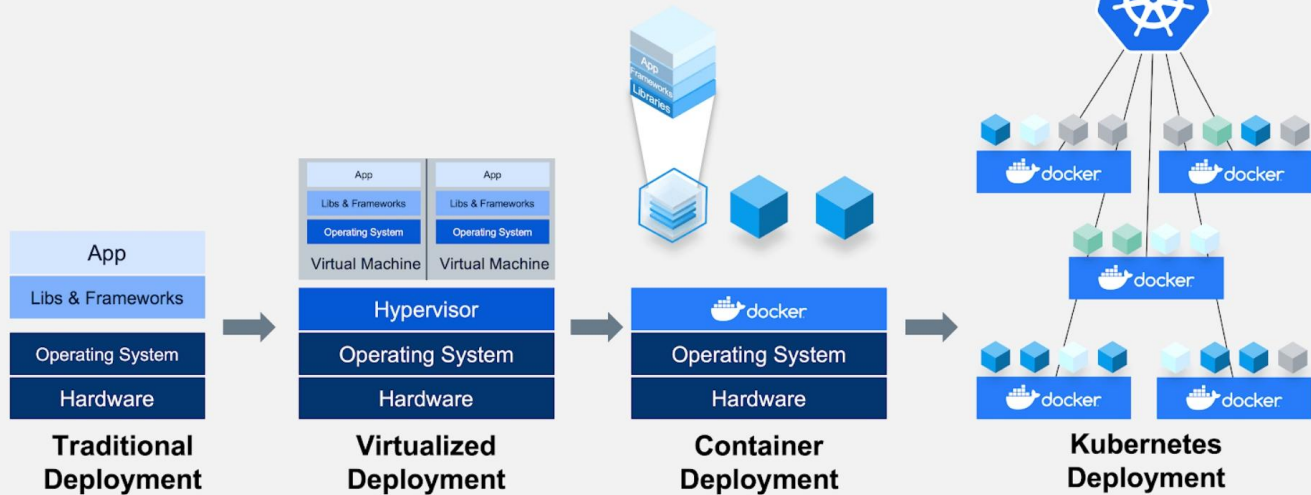
- What is container & Why?
- Container vs Virtual Machine
- Linux Containers & Docker
- Terminologies in Docker world
- Docker Architecture



# Container - Journey




## The past and the present of Apps Deployment



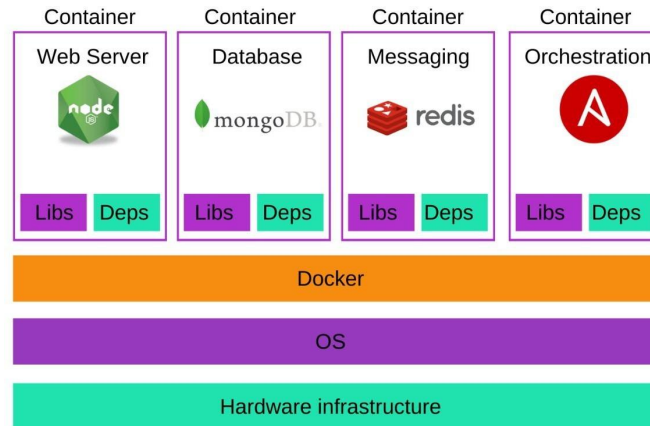
Kubernetes & Docker work together to build & run containerized applications

# What is Docker and Container ?

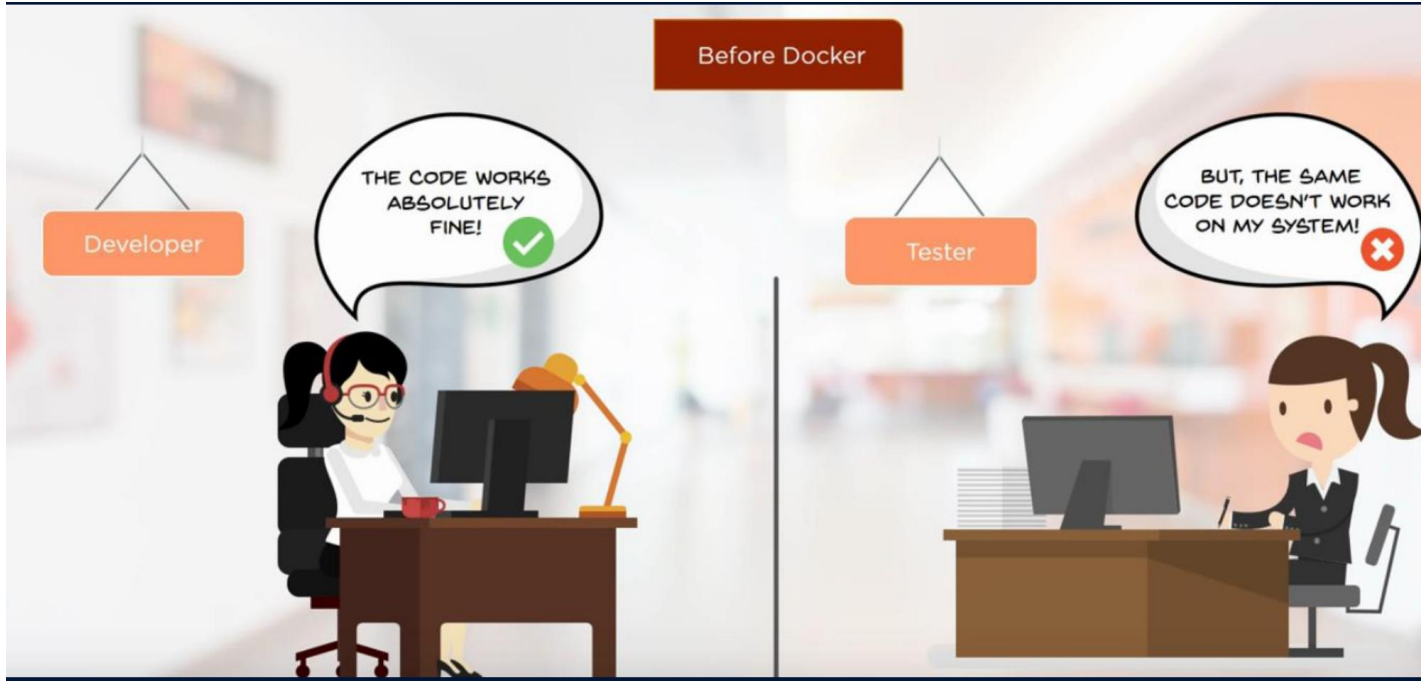


**Docker** is a software development tool and a virtualization technology that makes it easy to develop, deploy, and manage applications by using containers.

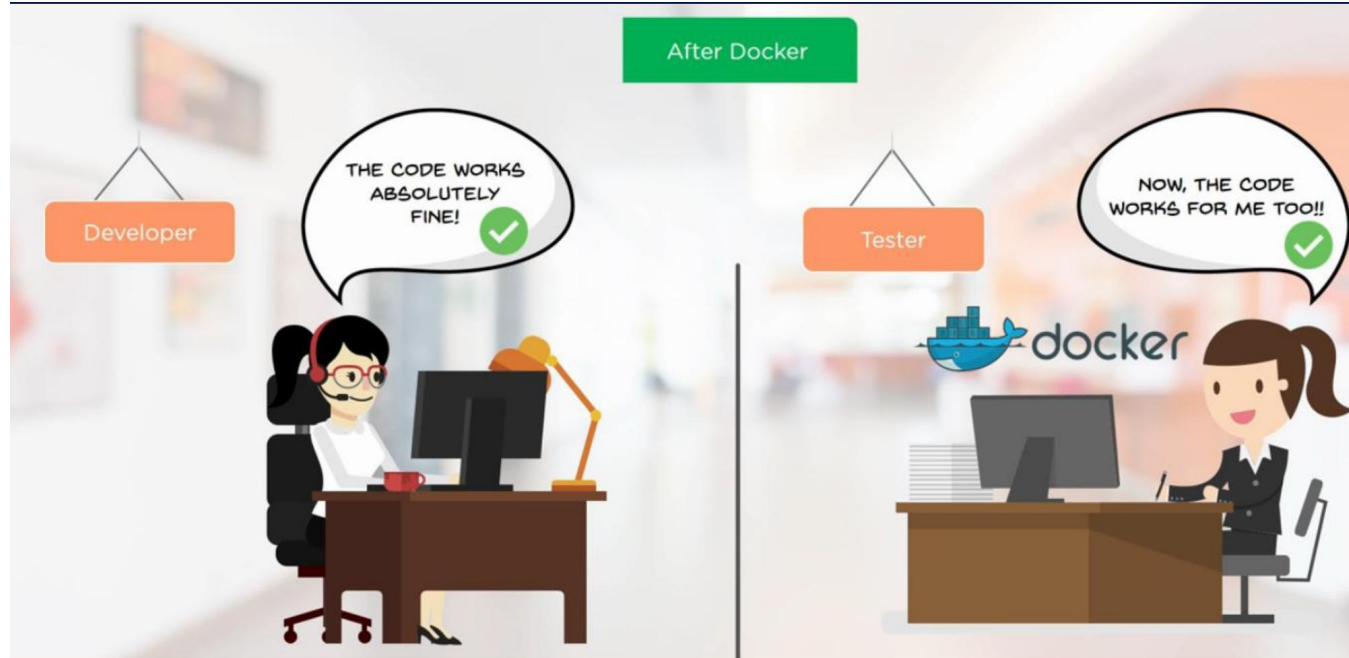
**Container** refers to a lightweight, stand-alone, executable package of a piece of software that contains all the libraries, configuration files, dependencies, and other necessary parts to operate the application.



# Why Containers?



# Why Containers?





# Why Containers?

---



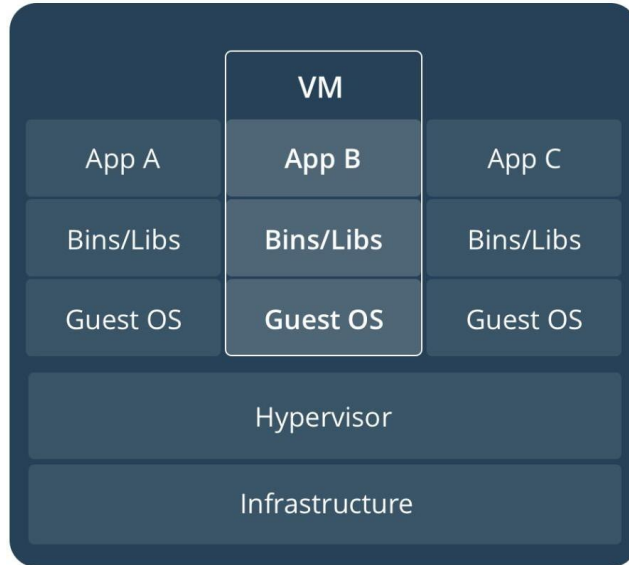
- **Flexible:** Even the most complex applications can be containerized.
- **Lightweight:** Containers leverage and share the host kernel.
- **Interchangeable:** You can deploy updates and upgrades on-the-fly.
- **Portable:** You can build locally, deploy to the cloud, and run anywhere.
- **Scalable:** You can increase and automatically distribute container replicas.
- Running more workload on the same hardware



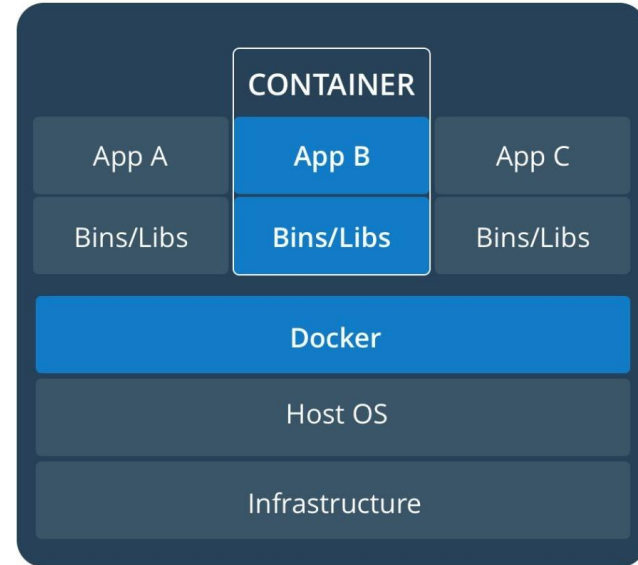
# Virtual Machines and Containers




Virtual Machine diagram



Container diagram



# Virtual Machines and Containers



Virtual Machine	Docker Container
Hardware-level process isolation	OS level process isolation
Each VM has a separate OS	Each container can share OS
Boots in minutes	Boots in seconds
VMs are of few GBs	Containers are lightweight (KBs/MBs)
Ready-made VMs are difficult to find	Pre-built docker containers are easily available
VMs can move to new host easily	Containers are destroyed and re-created rather than moving
Creating VM takes a relatively longer time	Containers can be created in seconds
More resource usage	Less resource usage

# Terminologies

---



**Image** Executable package that includes everything needed to run an application – the code, a runtime, libraries, environment variables, and configuration files

**Container** Runtime instance of an image—what the image becomes in memory when executed

**Service** a container but service codifies the way image runs -replicas, port, name etc

**Swarm** cluster of machines running docker containers

**Registry** storage and content delivery system, holding named Docker images, available in different tagged versions

**Server Daemon** creates and manages docker objects - images, containers, network, volumes, swarm etc

**Docker Client** CLI to communicate with server using Docker API

**Docker REST API** Communication contract between docker component (servers & clients)

**Network** Docker object holding the networking meta-data

**Node** machine participating in Swarm

**Volume** Storage of persistence data generated and managed by Docker containers




# Docker


## Architecture

---

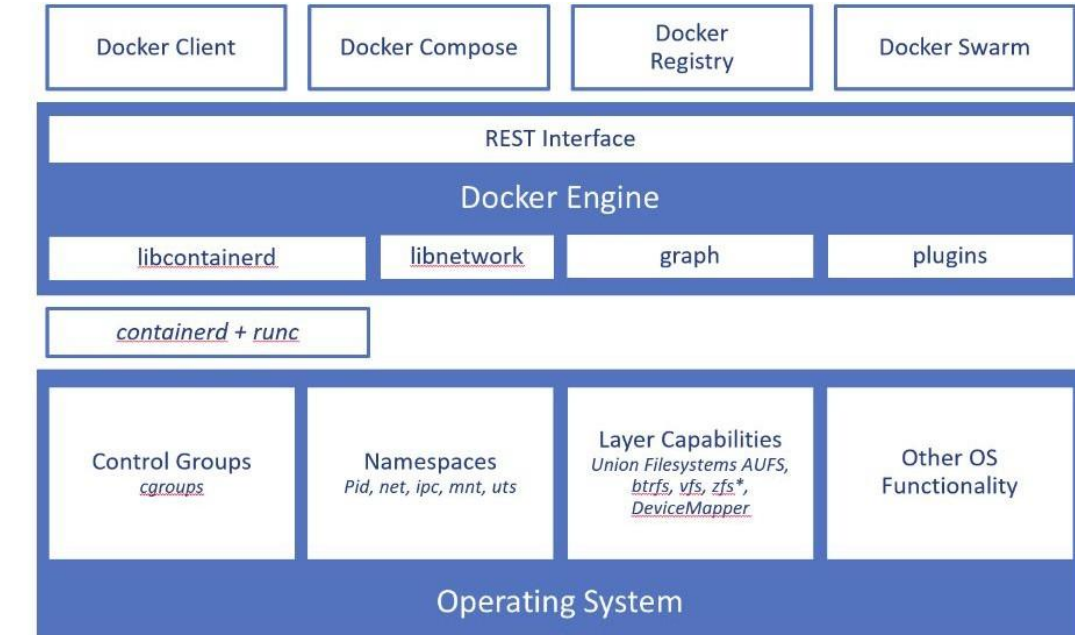
- Docker uses a client-server architecture.
- Docker client talks to the Docker daemon, which does the heavy lifting of building, running, and distributing your Docker containers.
- Docker client and daemon can run on the same system, or you can connect a Docker client to a remote Docker daemon.
- For a virtual communication between CLI client and Docker daemon, a REST API is used



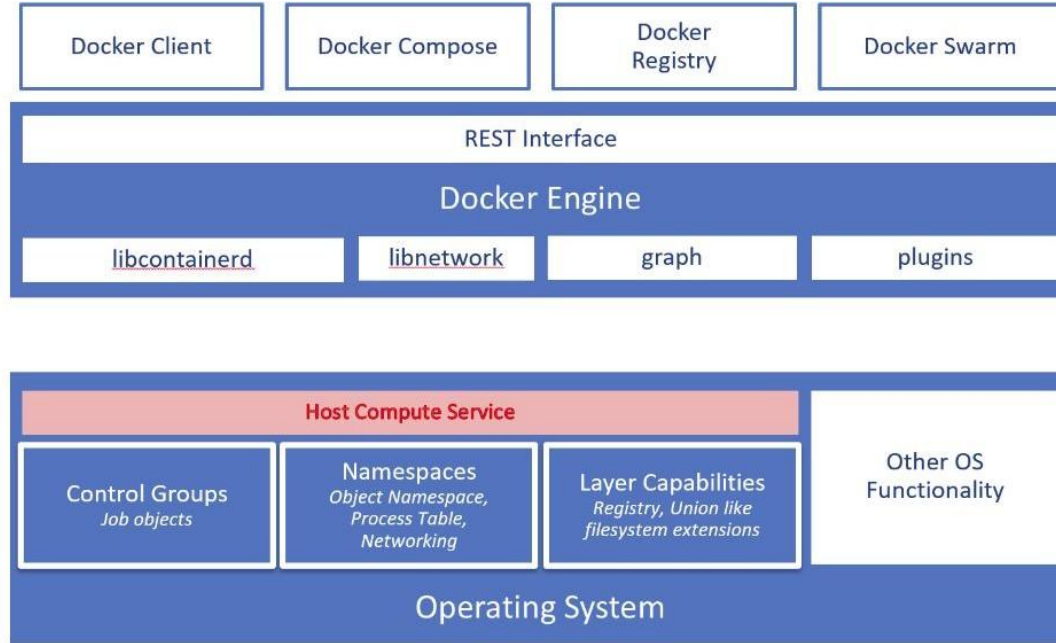
# Docker Architecture



# Docker Architecture - Linux



# Docker Architecture - Windows





# Docker Setup (Ubuntu)

---



```
sudo apt-get update
```

```
sudo apt-get remove docker docker-engine docker.io
```

```
sudo apt install docker.io
```

```
sudo groupadd docker
```

```
sudo usermod -aG docker $USER
```

```
sudo systemctl start docker
```

```
sudo systemctl enable docker
```



# Module 2: Docker Containers

---



- Creating & Starting containers
- Running containers
- Docker Images
- Connecting containers



# Creating containers

---



`docker container create [OPTIONS] IMAGE [COMMAND] [ARG...]`

Options:

- `--name` string name of the container
- `--cpus` decimal number of CPUs
- `--label` list set metadata on a container
- `--memory` bytes memory limit
- `--network` string connect container to a network (default “default”)
- `--publish` list publish container’s port to the host
- `--rm` remove container when it exits
- `-i` interactive Keep STDIN open if not attached
- `-t` allocates pseudo-TTY



# Creating containers - Examples



docker container create **--name hello-docker** alpine ping docker.com

command

argument(s)

options

image name from docker  
hub

docker container create --name busy -it busybox

docker container create --name alpine -it alpine sh

docker container create --name hello -p 80:80 tutum/hello-world

# Starting containers

---



`docker container start [OPTIONS] CONTAINER [CONTAINER...]`

## Options:


- i Attach container's STDIN
- a Attach container's STDOUT/STDERR and forward signals

## Examples:

`docker container start -ia busy`

`docker container start -ia alpine`

`docker container start hello`



# Running containers

---



`docker container run [OPTIONS] IMAGE [COMMAND] [ARG...]`

Options:

- `--name` string name of the container
- `--cpus` decimal number of CPUs
- `--label` list set metadata on a container
- `--memory` bytes memory limit
- `--network` string connect container to a network (default “default”)
- `--publish` list publish container’s port to the host
- `--rm` remove container when it exits
- `-i` interactive mode
- `-t` allocates a pseudo-TTY

# Running containers - Examples

---



`docker container run -p 80:80 tutum/hello-world` (creates container with random name)

`docker container run -p 80:80 nginx` (connects to tty, Ctrl+C to exit)

`docker run -p 81:80 nginx` (shorthand command)

`docker run --name ngx -p 80:80 -it nginx` (interactive terminal, Ctrl+PQ to leave it running)

`docker attach ngx`

`docker run -d ubuntu /bin/sh -c "while true; do echo current date and time is: $(date); sleep 10; done"`

`docker run -P --name nginx nginx` ( map exposed ports to random ports on the host {range 49153 and 65535})

`docker run -d -p 8000-9000:80 nginx` (maps port 80 to any random port between 8000 to 9000 on host)

`docker run --restart always -p 80:80 -it nginx`



# Docker Images

---



- **Image** - Executable package that includes everything needed to run an application – the code, a runtime, libraries, environment variables, and configuration files

**docker images**

**docker images nginx**

**docker images ubuntu**

**docker images ubuntu**

**docker rmi ubuntu**

**docker rmi \$(docker images -q)**





# Module 3: Provisioning Docker Images

---



- Introducing the Dockerfile
- Building images manually / Examples...
- Storing and retrieving Docker Images from Docker Hub
- Building images using Continuous Integration tools
- Inspecting a Dockerfile from DockerHub



# Different ways to create images

---

<b>docker commit</b>	Build an image from a container
<b>docker build</b>	Create an image from a Dockerfile by executing the build steps given in the file
<b>docker import</b>	Create a base image by importing from a tarball. [import is mainly used for creating base-images; first two options are widely used]

Dockerfile is a text document that contains

- a set of instructions required to assemble the app (image) and/ run it



# Introducing the Dockerfile



## Image Environment Blueprint

install node

set MONGO\_DB\_USERNAME=admin

set MONGO\_DB\_PWD=password

create /home/app folder

copy current folder files to /home/app

start the app with: "node server.js"

## DOCKERFILE

**FROM** node

**ENV** MONGO\_DB\_USERNAME=admin \  
MONGO\_DB\_PWD=password

**RUN** mkdir -p /home/app

**COPY** . /home/app

**CMD** ["node", "server.js"]

blueprint for building images

# Introducing the Dockerfile

---

- **ENV** - to set environment variables
- **EXPOSE** - to expose ports
- **FROM** - base image
- **LABEL** - to add metadata to image
- **HEALTHCHECK** - to check if container is running
- **USER** - to set user and group
- **VOLUME** - to specify mount point from external host
- **WORKDIR** - workdir to run any of the commands



# Introducing the Dockerfile

---

- ARG - variable used during build time
- CMD - to provide defaults to executing container
- RUN - to execute commands in new layer
- COPY - Copy file,dir or remote url to image
- ADD - Copy file,dir or remote url to image
- ENTRYPOINT - to configure container as executable
- MAINTAINER - the image maintainer

RUN COPY ADD instructions create new layers in the image stack - refer layering section



# Building Images (Alpine ping)

---



cat Dockerfile

```
FROM alpine:latest  
MAINTAINER ninad@gmail.com  
CMD ping google.com
```

## Build

- `docker build -t myalpine .`

## Run

- `docker run myalpine`



# Building Images (Ubuntu with utilities)

---



cat Dockerfile

FROM ubuntu:latest

MAINTAINER [ninad@gmail.com](mailto:ninad@gmail.com)

RUN apt-get update && apt-get install -y tree && apt-get install -y telnet && apt-get install -y curl

Build

- `docker build -t myubuntu`
- `.`

Run

- `docker run -it myubuntu`



CGI

# Building Images (hello-world)

---

```
git clone https://github.com/NinjaCloud/tutum-hello-world-rebuild
```

```
FROM nginx:1.19.6-alpine
```

```
RUN apk --update add php-fpm
```

```
RUN mkdir -p /tmp/nginx && echo "clear_env = no" >> /etc/php7/php-fpm.conf
```

```
ADD www /www
```

```
ADD nginx.conf /etc/nginx/
```

```
COPY entrypoint.sh .
```

```
RUN chmod +x entrypoint.sh
```

```
ENTRYPOINT [ "./entrypoint.sh" ]
```

## Build

- `docker build -t hello-world .`

## Run

- `docker run -it hello-world`



# Building Images (Java-Program)



cat Dockerfile

```
FROM java:latest
COPY . /usr/src/
WORKDIR /usr/src/
RUN javac hello.java
CMD ["java", "hello"]
```

vi hello.java

```
class hello {
    public static void main(String []args) {
        System.out.println("Hey Ninad");
    }
}
```

Build

- `docker build -t mycode-java .`

Run

- `docker run -it mycode-java`

# Docker Hub - store & retrieve

---



<https://hub.docker.com> (register and create login)

- `docker tag alpine ninjacloud05/alpine:ninad`
- `docker push ninjacloud05/alpine:ninad`
- `docker pull ninjacloud05/alpine:ninad`

```
docker login -u "myusername" -p "mypassword" docker.io  
docker push myusername/myimage:0.0.1
```



# Module 4: Diving deeper - Dockerfile

---

- Dockerfile and Layers
- The Build cache
- The ENTRYPOINT Instruction
- The CMD Instruction Docker
- The ENV Instruction
- Volumes and the VOLUME Instruction



# Dockerfile & Layers



```
ubuntu@ip-172-31-31-236:~$ docker images springio/*
REPOSITORY TAG IMAGE ID CREATED SIZE
springio/gs-spring-boot-docker latest 3a7a85f42b64 6 months ago 181MB

ubuntu@ip-172-31-31-236:~$ docker history 3a7a85f42b64
IMAGE CREATED CREATED BY SIZE COMMENT
3a7a85f42b64 6 months ago /bin/sh -c #(nop) ENTRYPOINT ["sh" "-c" "... 0B
<missing> 6 months ago /bin/sh -c #(nop) ENV JAVA_OPTS= 0B
<missing> 6 months ago /bin/sh -c #(nop) ADD file:2f6c6463d5fd2c4... 14.4MB
<missing> 6 months ago /bin/sh -c #(nop) VOLUME [/tmp] 0B
<missing> 6 months ago /bin/sh -c apk add --no-cache --virtual=bu... 156MB
<missing> 6 months ago /bin/sh -c #(nop) ENV JAVA_VERSION=8 JAVA... 0B
<missing> 7 months ago /bin/sh -c #(nop) ENV LANG=C.UTF-8 0B
<missing> 7 months ago /bin/sh -c ALPINE_GLIBC_BASE_URL="https://... 6.7MB
<missing> 7 months ago /bin/sh -c #(nop) CMD ["/bin/sh"] 0B
<missing> 7 months ago /bin/sh -c #(nop) ADD file:4583e12bf5caec4... 3.97MB
```

# Build Cache

---



## Why Layers & Cache?

- To identify similar portions of content by componentizing image
- To avoid downloading similar content thus reduce network traffic
- To build images faster by reusing parts which were created earlier



# ENTRYPOINT and CMD



FROM Ubuntu

CMD sleep 5

CMD command param1

CMD ["command", "param1"]

CMD sleep 5

CMD ["sleep", "5"]



CMD ["sleep 5"]



```
docker build -t ubuntu-sleeper .
```

```
docker run ubuntu-sleeper
```



5

# ENTRYPOINT and CMD



FROM Ubuntu

CMD sleep 5

```
docker run ubuntu-sleeper sleep 10
```

Command at Startup: sleep 10

FROM Ubuntu

ENTRYPOINT ["sleep"]

```
docker run ubuntu-sleeper 10
```

Command at Startup: sleep 10

```
docker run ubuntu-sleeper
sleep: missing operand
Try 'sleep --help' for more information.
```

Command at Startup: sleep

# ENTRYPOINT and CMD



FROM Ubuntu

ENTRYPOINT ["sleep"]

CMD ["5"]

```
▶ docker run ubuntu-sleeper  
sleep: missing operand  
Try 'sleep --help' for more information.
```

Command at Startup: sleep 5

```
▶ docker run ubuntu-sleeper 10
```

Command at Startup: sleep 10

```
▶ docker run --entrypoint sleep2.0 ubuntu-sleeper 10
```

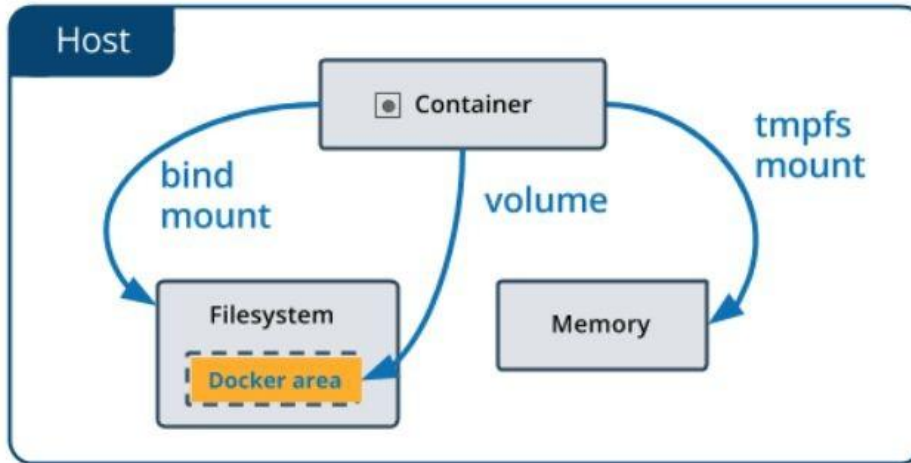
Command at Startup: sleep2.0 10



# The VOLUME - Data Persistence



Storage of persistence data generated by managed by Docker containers



Commands:

- `docker volume create my-vol`
- `docker volume ls`
- `docker volume inspect my-vol`
- `docker volume rm my-vol`

# VOLUME - Examples

---



Examples (volume): Persist data in a container's writeable layer

- `docker run -d --name devtest --mount source=app,target=/app nginx:latest`

Examples (bind volume): a file or directory on the *host machine* is mounted into a container. Performant but not-reliable

- `docker run -d -it --name devtest --mount type=bind,source="$(pwd)",target=/app nginx:latest`

Examples (tmpfs volume): For temporary sensitive data to be kept only in memory

- `docker run -d -it --name tmpptest --mount type=tmpfs,destination=/app nginx:latest`



# VOLUME - preferred way

---



- Volumes are easier to back up or migrate than bind mounts.
- You can manage volumes using Docker CLI commands or the Docker API.
- Volumes work on both Linux and Windows containers.
- Volumes can be more safely shared among multiple containers.
- Volume drivers allow you to store volumes on remote hosts or cloud providers, to encrypt the contents of volumes, or to add other functionality.
- A new volume's contents can be pre-populated by a container.



# Module 5: Working with Registry

---

- Overview
- Creating a Public repo on Docker Hub
- Using our Public repo on Docker Hub
- Using a Private Registry
- Docker Enterprise
- Lab Exercises

# Overview - Registry

---



## Registry

Stateless, highly scalable server side application that stores and lets you distribute Docker images.

## When to use

- tightly control where your images are being stored
- fully own your images distribution pipeline
- integrate image storage and distribution tightly into your in-house development workflow



# Registry Server

---





## Run a local registry

```
$ docker run -d -p 5000:5000 --restart=always --name registry registry:2
$ docker pull ubuntu:16.04
$ docker tag ubuntu:16.04 localhost:5000/my-ubuntu
$ docker push localhost:5000/my-ubuntu
$ docker image remove ubuntu:16.04
$ docker image remove localhost:5000/my-ubuntu
$ docker pull localhost:5000/my-ubuntu
$ docker container stop registry && docker container rm -v registry
```




# Jenkins

---



# Dockerizing dev workflow



```
docker run --name jenkins -u 0 -d -p 8080:8080 -v /var/run/docker.sock:/var/run/docker.sock -v $(which docker):$(which docker) jenkins/jenkins:lts
```

## Notes:

Add docker pipeline jenkins plugin to work

Test project: <https://github.com/NinjaCloud/nodejsappdocker.git>

Add jenkins credential having ID **docker-hub-credentials** for docker hub push access



# Module 6: Docker Networking

---



- Overview
- The docker0 Bridge
- User Defined Network
- Exposing Ports
- Viewing Exposed Ports
- Linking Containers



# Overview - Networking

---



Defines how containers communicate with external world, amongst cluster members etc

Two types of networks:

- Default
- Custom Defined

Default:

- Bridge - docker0 (docker created default network) **Configurable**
- Host - container on host network stack **Not configurable**
- None - container specific network stack (no network interface) **Not configurable**



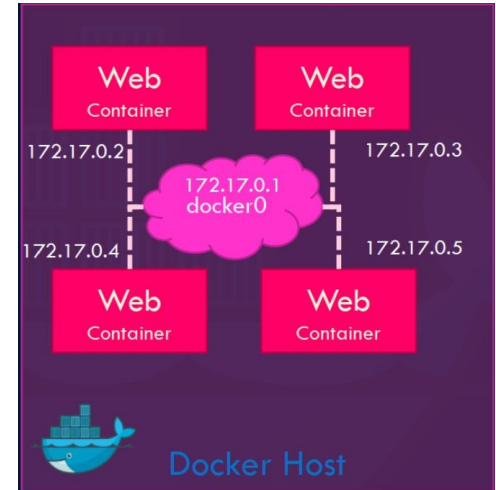
# The docker0 bridge



- In Bridge network, all containers get private internal IPs and they are isolated from host.
- Container inter-connectivity using IP addresses (no name resolution)
- For name resolution, legacy `--link` feature available for limited period
- Change default bridge to none using `--network` flag or `daemon.json` server config

Create a bridge network: `docker network create --driver bridge my-net`

Attach a container to it: `docker run -d --name web --net my-net nginx`



# Understanding DNS resolution in bridge network

- When containers are run in default bridge network they cannot find each other using their container names.
- Simply put, DNS resolution through container names will not work under default bridge network

```
ninad@ninad-X555LAB: ~  
ninad@ninad-X555LAB: ~ 150x38  
ninad@ninad-X555LAB:~$ docker run -d --name c1 alpine sleep 3600  
50021b776f9d2e674ce9dc01d488782a13907120b65f0367e111423e98a3f059  
ninad@ninad-X555LAB:~$ docker run -d --name c2 alpine sleep 3600  
b78fcc47f22b865c5be230b65d967533581e28b336aba3c202677adcb0985209  
ninad@ninad-X555LAB:~$ docker ps  
CONTAINER ID   IMAGE     COMMAND                  CREATED    STATUS    PORTS                               NAMES  
b78fcc47f22b   alpine    "sleep 3600"            8 seconds ago    Up 6 seconds      
50021b776f9d   alpine    "sleep 3600"            14 seconds ago    Up 13 seconds      
b41d82b8ae97   nginx     "/docker-entrypoint..." 21 hours ago     Up 21 hours     0.0.0.0:8080->80/tcp, :::8080->80/tcp    sharp_spence  
ef6e6758b080   registry:2 "/entrypoint.sh /etc..." 5 months ago     Up 23 hours     0.0.0.0:5000->5000/tcp, :::5000->5000/tcp registry  
ninad@ninad-X555LAB:~$ docker exec -it c1 bash  
OCI runtime exec failed: exec failed: container_linux.go:380: starting container process caused: exec: "bash": executable file not found in $PATH: unknown  
ninad@ninad-X555LAB:~$ docker exec -it c1 sh  
/ # ping c2  
ping: bad address 'c2'  
/ #
```

# Understanding DNS resolution in bridge network

- Now a new bridge network is created and containers are attached to that network.
- In this case, containers find each other using their container names(DNS resolution through container names)


```
ninad@ninad-X555LAB: ~  
ninad@ninad-X555LAB: ~ 150x38  
ninad@ninad-X555LAB:~$ docker network create -d bridge mynet01  
24ba77fa573c818702804affe45daa5bb152150d31e8ce2a413a5c414128bb28  
ninad@ninad-X555LAB:~$ docker run -d -it --net mynet01 --name c1 alpine sleep 3600  
874c6eb18c29159cd45c41d40435655ba7a0258020467ae3216e04e6784b55fd  
ninad@ninad-X555LAB:~$ docker run -d -it --net mynet01 --name c2 alpine sleep 3600  
baece5e6467cb573e3e40703b239e97a458c54277ae4d9f40823c6c0b0c36638  
ninad@ninad-X555LAB:~$ docker ps  
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                               NAMES  
baece5e6467c   alpine    "sleep 3600"            24 seconds ago Up 22 seconds                               c2  
874c6eb18c29   alpine    "sleep 3600"            33 seconds ago Up 31 seconds                               c1  
b41d82b8ae97   nginx     "/docker-entrypoint...." 21 hours ago   Up 21 hours   0.0.0.0:8080->80/tcp, :::8080->80/tcp sharp_spence  
ef6e6758b080   registry:2 "/entrypoint.sh /etc..." 5 months ago   Up 23 hours   0.0.0.0:5000->5000/tcp, :::5000->5000/tcp registry  
ninad@ninad-X555LAB:~$ docker exec -it c1 sh  
/ # ping c2  
PING c2 (172.19.0.3): 56 data bytes  
64 bytes from 172.19.0.3: seq=0 ttl=64 time=0.278 ms  
64 bytes from 172.19.0.3: seq=1 ttl=64 time=0.220 ms  
64 bytes from 172.19.0.3: seq=2 ttl=64 time=0.189 ms  
64 bytes from 172.19.0.3: seq=3 ttl=64 time=0.212 ms  
64 bytes from 172.19.0.3: seq=4 ttl=64 time=0.188 ms  
64 bytes from 172.19.0.3: seq=5 ttl=64 time=0.191 ms
```

# Docker Networking: Host



- In host network, all containers directly get connected to host.
- Multiple containers cannot run on same hosts because of port conflicts on host side

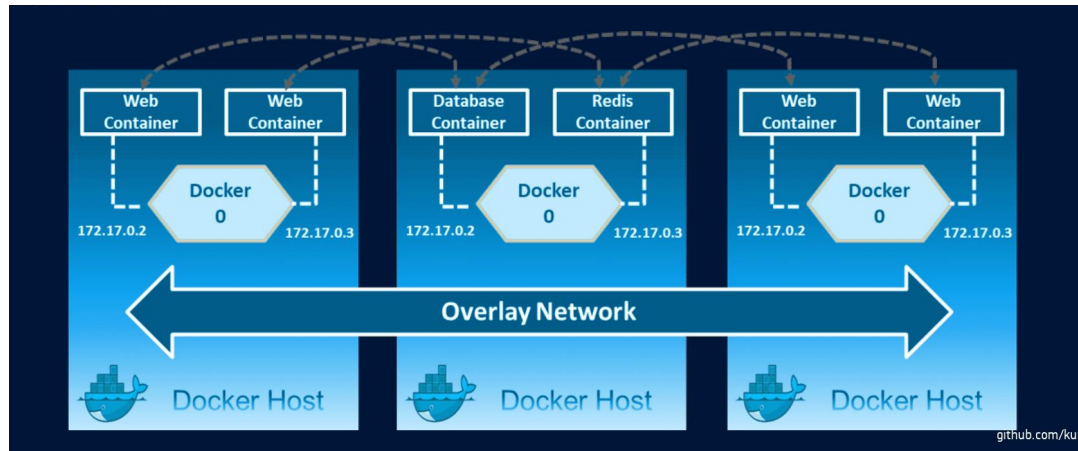
`docker run -d --name web --net host nginx`



# Overlay Network



- Scope is swarm mode
- Bridge networks apply to containers running on the same Docker daemon host. For communication among containers running on different Docker daemon hosts, we should use an overlay network which spans across the entire cluster
- Uses NAT and port mapping (iptables)



# Test Setup - 1

---



## Test Setup:

Create custom network n1

- `docker network create n1`

Create two busybox containers attached to n1

- `docker run -itd --name c1 --network n1 busybox`
- `docker run -itd --name c2 --network n1 busybox`

## Tests

- Log into c1 and ping c2 (should succeed)
  - `docker exec -it c1 sh`
  - `ping c2`
- Log into c2 and ping c1 (should succeed)
  - `docker exec -it c2 sh`
  - `ping c1`



# Test Setup - 2

---



**Prerequisites:** Test Setup -1

**Test Setup:**

Remove network from both containers c1 & c2

- `docker network disconnect n1 c1`
- `docker network disconnect n1 c2`

**Tests:**

- Login into c1 and ping c2 (should fail)
  - `docker exec -it c1 sh`
  - `ping c2`
- Login into c1 and ping google.com (should fail)
  - `docker exec -it c1 sh`
  - `ping google.com`
- Run `ifconfig` on c1 to see interfaces (should see only loopback interface)
  - `docker exec -it c1 sh`
  - `ifconfig`
- Do the same on c2 (results should be similar)



# THANK YOU !!

